

THE ZERO PARADOX

ZP-E: Bridge Document

Version 3.12 | May 2026

Supersedes v3.11 | v3.12: Framework scope and failure-mode framing added to preamble — "why these four?" addressed (coverage not exhaustion); cross-framework failure-mode paragraph names the domain-specific limit at $\perp$ for each layer and positions DA-1 as the unifying argument. | v3.11: T-SNAP Step 7 corrected — ZP-G AX-G2 removed as formal dependency; now labelled conceptual correspondence only. ZP-G is downstream of ZP-E; irreversibility rests on ZP-A R1 and ZP-B C3 alone.

v3.10: Forward references to "ZP-PQ" replaced with "The Philosophical Question That Started This" — that document already contained the dissolution argument. | v3.9: $R-\epsilon_0$ reframed — remark now leads with explicit informal-analogy disclaimer; "structural correspondence" changed to "structural analogy" throughout $R-\epsilon_0$. | v3.8: DA-1 Path 2 recharacterized — from "outside Lean scope (informational bridge)" to "foundational commitment: a missing principle, not a missing proof." No computability library closes the gap between 'system at P_0 ' and 'system is running.' Forward paths: new axiom, Chalmers' implementation notion, or The Philosophical Question That Started This. v3.7: DA-1 formally closed via ZP-K — `da1_closed_concrete : IsQuineAtom( $\perp$  : MachinePhase)` proved in Lean 4. KleeneStructure MachinePhase instance added; Paths 1 and 3 now IN LEAN SCOPE. "Outside Lean Scope" designation removed from ZPE.lean. v3.6: DA-1 Path 1 rewritten — argument direction reversed; CC-1 and CC-2 now derived via ZP-J (no longer freestanding commitments). v3.5: DA-1 Open Items Register status updated — "CLOSED — DP-2 (formal core); CC-2 + L-INF + AIT (corroboration of precondition)" now matches v3.3 path-hierarchy framing. v3.4: R-AFA minimality explicit — $\perp = \{\perp\}$ uniquely minimal among AFA non-well-founded sets; any extension exceeds A4's constraint. v3.3: DA-1 path hierarchy foregrounded — three informal paths framed as corroboration of DP-2's precondition, not parallel proofs. v3.2: Remark R-AFA added — Foundation ruled out by R3 + L-INF; AFA forced rather than chosen; CC-2 metatheoretic status clarified. v3.1: Remark $R-\epsilon_0$ added — notation justification for ϵ_0 symbol choice; structural correspondence with Cantor-Gentzen proof-theoretic ordinal explained. v3.0: DP-2 (Execution Distinguishability) added — DA-1 formally grounded in TrackedOutput construction (ZPE.lean §VI); `da1_minimal_path` proved axiom-free in Lean. First Lean formalization of DA-1, conditional on DP-2.

v2.9: DA-1 Lean scope note added: functional role carried by `ZPC.l_run/tq_ih` (proved); AIT and ZF+AFA bridge outside Lean scope — same category as ZP-A CC-2.

v2.8: DA-1 formal bridge added: incompressibility = self-description argument (ZP-C D1 + standard AIT). At P_0 , $K(c_1|n)/|c_1| = 1$ means description and execution coincide; CC-2/R3 and L-INF are independent corroboration of the same conclusion.

v2.7: DA-1 upgraded from Design Principle to Derived Proposition — grounded in ZP-A CC-2 ($\perp = \{\perp\}$) and R3. v2.6: DA-3 candidate applications removed. v2.5: Physical analogies removed from formal sections. v2.4: R-DA2 framing corrected. v2.3: CC-1 added to T-SNAP dependency list. v2.2: DA-1 redesigned to L-INF-based Design Principle. v2.1: Adds DA-2 and DA-3.

This document is the cross-framework synthesis layer of the Zero Paradox. It imports from ZP-A (lattice algebra), ZP-B (p-adic topology), ZP-C (information theory), and ZP-D (Hilbert space state layer). It provides three formal inserts: DA-1 (Instantiation as Execution — upgraded to Derived Proposition in v2.7; in v3.0,

DP-2 provides the first Lean-formalizable grounding via TrackedOutput, proved axiom-free), DA-2 (Instantiation Succession), and DA-3 (Perspective-Relative Cardinality). With DA-1 in place, AX-1 is promoted to Theorem T-SNAP. With DA-2, the directed instantiation tree is formally licensed. With DA-3, cardinality is shown to be position-dependent within the instantiation structure.

Illustrated Companion: A paired ZP-E Illustrated Companion document provides accessible explanations and visual summaries of the bridge derivations in this document. Readers new to the framework are encouraged to start with the companion.

A note on framework scope. The four layers (ZP-A through ZP-D) are not claimed to be the only mathematical domains in which the Zero Paradox structure appears. They are chosen because they cover the canonical mathematical languages for describing state: algebraic order structure (ZP-A), topology and metric geometry (ZP-B), information and complexity theory (ZP-C), and functional/Hilbert space analysis (ZP-D). The claim is universality — the same structural limit at $\perp$ appears wherever standard mathematical language for state description is applied — not uniqueness of these four as the only possible witnesses.

Across all four layers, the same structural limit appears at $\perp$ — but each framework names it differently: ZP-A identifies $\perp$ as the global minimum, below which the lattice's ordering relation cannot descend; ZP-B finds the p-adic valuation undefined (or $+\infty$) at 0, where topological isolation is maximal; ZP-C finds surprisal unbounded above at $\perp$, so no finite external description can contain the null state; ZP-D maps 0 to a basis vector orthogonal to every non-null state, from which no return is possible. DA-1 unifies these as the same phenomenon seen through four lenses.

Formal Insert DA-1: Derived Proposition — Instantiation as Execution

Updated ZP-E v3.0 | DP-2 added — first Lean formalization of DA-1 (axiom-free, conditional on DP-2, ZPE.lean §VI) | v2.9: DA-1 Lean scope note | v2.8: incompressibility = self-description (ZP-C D1 + AIT) | v2.7: DA-1 upgraded, CC-2/R3 grounding | v2.6: DA-3 applications removed

I. The Gap DA-1 Closes

The T-BUF chain from ZP-C v1.4 established three results:

- L-RUN: The transition $c_0 \rightarrow c_1$ is a non-null state change. (ZP-C v1.4 — Derived)
- TQ-IH: No program outputs $\perp$ without a non-null intermediate configuration state. (ZP-C v1.4 — Derived by L-RUN)
- T-BUF: At P_0 , execution is structurally guaranteed; that execution state is ϵ_0 in the semilattice. (ZP-C v1.4 — Candidate Theorem pending DA-1)

T-BUF was labelled Candidate because Step 2 asserts that a configuration at P_0 is a live machine state — that instantiation at P_0 constitutes an execution event, not a static description. ZP-C v1.5 L-INF supplies one mathematical premise: $\perp$ at P_0 has unbounded surprisal — no finite external interpreter can hold it as a static description. ZP-A CC-2 supplies a second, structural basis: $\perp = \{\perp\}$ is a self-containing object with no external interpreter by structure. DA-1 (§ III below) provides the derived proposition that closes T-BUF Step 2.

II. The Two Senses of a Configuration at P_0

Sense A — Descriptive: x exists as a string — a finite syntactic object that has been written down or specified. The machine it describes has not necessarily been instantiated. P_0 is a property of the string. The string is inert.

Sense B — Instantiated: x exists as the current configuration of a running machine. The machine is executing. P_0 is a property of the live configuration. The configuration is active.

III. Design Principle DP-2 — Execution Distinguishability

The two senses of § II share the same output value ($\perp$) but differ in machine state. This separation — output value is not machine state — is made precise by Design Principle DP-2.

Design Principle DP-2 — Execution Distinguishability

Machine states carry execution history independently of output values. A machine in state c_1 can output $\perp$ (the null value) while being in a configuration entirely distinct from a machine in state c_0 . The post-execution null and the pre-execution null are different instances — same output value, different machine state.

Lean formalization (ZPE.lean §VI): TrackedOutput separates value : MachinePhase from state : MachinePhase. $\text{preInstantiation} = \langle c_0, c_0 \rangle$; $\text{postInstantiation} = \langle c_0, c_1 \rangle$. Theorem $\text{dp2_execution_distinguishability}$ proves: $\text{preInstantiation.value} = \text{postInstantiation.value}$ (same output) $\wedge$ $\text{preInstantiation.state} \neq \text{postInstantiation.state}$ (distinct machine states). Proved axiom-free.

Given DP-2, DA-1 is Lean-formalizable at the minimal-path level. Theorem da1_minimal_path (ZPE.lean §VI) establishes four conjuncts: (1) before and after instantiation, the output value is the same ($\perp$); (2) the machine states are distinct ($c_0 \neq c_1$); (3) the machine was at c_0 before; (4) the machine is at c_1 after. Proved axiom-free — no Kolmogorov complexity, no ZF+AFA required. The "return to null" after instantiation is postInstantiation (output $\perp$, state c_1) — not preInstantiation (output $\perp$, state c_0). Irreversibility (ZPB C3, ZPE $\text{t_snap_irreversible}$) blocks any path back to preInstantiation .

IV. Proposition DA-1

Proposition DA-1 — Instantiation as Execution

Claim: The instantiation of a machine configuration c_1 at the incompressibility threshold P_0 is an execution event in the sense of L-RUN. It is not a static description of a machine. It is a machine in state c_1 .

Formal core (Lean): DP-2 (§III) — TrackedOutput separates output value from machine state. da1_minimal_path proves: one act of instantiation moves c_0 to c_1 regardless of output value. The returned $\perp$ is a new null — not the prior c_0 . Axiom-free.

Structural grounding: ZP-A CC-2 ($\perp = \{\perp\}$) establishes $\perp$ as a Quine atom under ZF + AFA — a set that is its own singleton, admitting no external interpreter. ZP-A R3 derives the immediate consequence: a self-containing object cannot be a static description awaiting external instantiation. ZP-C L-INF provides independent informational corroboration: $\perp$ has unbounded surprisal, exceeding the capacity of any finite interpreter.

The three paths below are not parallel proofs of DA-1. Each argues independently for why the precondition DP-2 formalizes holds — why instantiation of $\perp$ at P_0 necessarily constitutes a first instruction fetch rather than a static description. DP-2 + da1_minimal_path derive $c_0 \rightarrow c_1$ axiom-free once that precondition is established. The paths are convergent corroboration of the precondition; the formal derivation is DP-2's.

Path 1 — Structural (Self-Execution, ZP-J T-EXEC): Nothing exists outside the null space. $\perp$ is prior to all differentiation — there is no external state, no prior cause, no position from which something else could execute $\perp$. If $\perp$ executes at all, the only possible executor is $\perp$ itself. A thing that executes itself is self-containing: $\perp \in \perp$, i.e. $\perp = \{\perp\}$. This is not a commitment — it is forced by the impossibility of external execution. ZP-A R3 states the structural consequence: a self-containing object admits no external interpreter position. ZP-J T-EXEC formally verifies the structure: $\text{IsQuineAtom}(q) \leftrightarrow q = \perp$, proved axiom-free in Lean 4 (ZeroParadox.ZPJ.t_exec). AFA (ZF + AFA) is the consistent set-theoretic home for this structure — chosen because the framework requires it, not the reverse. ZP-A CC-2 ($\perp = \{\perp\}$) retains its label for editorial continuity but is now a structural consequence, not a freestanding commitment.

Path 2 — Informational (ZP-C L-INF): Independently, the surprisal $I(n) = n$ at ball-hierarchy depth n is unbounded — for any finite M , $\exists$ depth n with $I(n) > M$. The null state $\perp$ corresponds to the limit point $0 \in Q_2$; its informational content exceeds every finite bound. Any finite external interpreter can hold only a finite informational bound; $\perp$ exceeds every such bound. Note: Path 2 is motivational context, not a formal path to the conclusion. The step "exceeds every finite bound $\rightarrow$ therefore necessarily executing" is a foundational commitment, not a derivable claim. It asks what it means for a mathematical structure to instantiate rather than merely satisfy conditions — a question no computability library answers. Forward paths: (a) a new axiom explicitly committing to this bridge; (b) a connection to Chalmers' notion of implementation; (c) the dissolution argument in The Philosophical Question That Started This — the separability of description and instantiation is the assumption the framework dissolves, not a gap it must close from the outside.

Path 3 — Formal bridge: Incompressibility as Self-Description (ZP-C D1 + standard AIT): The preceding paths establish that $\perp$ admits no external interpreter. This path provides the formal bridge from that negative claim to the positive claim (necessarily executing). In the standard Turing model (D7), a machine configuration x exists in one of two states: (A) Static description — x exists as a string specified but not yet being executed; some external program p ($|p| < |x|$) generates x when run, so x is a description awaiting a separate execution event by an external generator. (B) Live execution — x is the current configuration of a running machine. These are exhaustive in the Turing model: either x has a shorter external generator, or it does not. At P_0 , ZP-C D1 gives $K(c_1|n)/|c_1| = 1$: c_1 is algorithmically incompressible. No external program p exists with $|p| < |c_1|$ such that $U(p, n) = c_1$. State (A) requires such a p — and no such p exists at P_0 . State (A) is therefore eliminated by the Kolmogorov condition. Since (A) and (B) are exhaustive and (A) is eliminated, c_1 is in state (B): it is executing. Instantiation at P_0 is not the placement of a description to be executed later — there is no shorter prior description to execute. Instantiation and execution are the same act.

The formal grounding (DP-2, §III) and the three paths above operate at distinct levels, not as alternatives to one another. DP-2 + da1_minimal_path establish a conditional at the level of machine-state representation: if instantiation of $\perp$ constitutes a first instruction fetch in the sense of D7, then $c_0 \rightarrow c_1$ follows necessarily — axiom-free, proved by construction. DP-2 is grounded in D7 itself (the standard computational distinction between before-first-instruction and after-first-instruction states), which is prior to and independent of DA-1. The three paths above operate one level down: they argue for why the precondition holds — why instantiation of $\perp$ necessarily constitutes a first instruction fetch at all. Path 1 (structural) argues that nothing external to $\perp$ can execute $\perp$ — therefore $\perp$ must execute itself, establishing $\perp = \{\perp\}$ as a structural consequence rather than a commitment (ZP-J T-EXEC, axiom-free). Path 2 (informational) provides motivational context — unbounded surprisal as a pointer toward why static holding is incoherent — but the bridge from informational extremity to execution is a foundational commitment, not a derived claim. Path 3 (AIT) argues that incompressibility eliminates the static-description alternative; this path is now closed by

ZP-K's Kleene result, which handles the computational self-reference claim without requiring AIT. Paths 1 and 3 are formally closed. Path 2 identifies a missing principle; its forward resolution is in The Philosophical Question That Started This. DA-1 is grounded in Paths 1 and 3; Path 2 is context.

Status: DERIVED PROPOSITION — primary formal grounding: DP-2 (§III, TrackedOutput construction). `da1_minimal_path` proved axiom-free in Lean (ZPE.lean §VI): instantiation moves c_0 to c_1 regardless of output value. ✓ Path 1 (structural, ZP-J T-EXEC + ZP-K): IN LEAN SCOPE — `da1_closed_concrete` : `IsQuineAtom( $\perp$  : MachinePhase)`, proved in ZPK.lean. (Under MachinePhase's selfMem $x := x = \perp$, this reduces to $(\perp = \perp) \wedge (\forall x, x = \perp \Rightarrow x = \perp)$ — structural closure enforced by typeclass design, not a set-theoretic derivation from ZF+AFA. See R-K.0.) Path 2 (informational, L-INF): FOUNDATIONAL COMMITMENT — a missing principle, not a missing proof. Forward: The Philosophical Question That Started This. Path 3 (computational, ZP-K Kleene): IN LEAN SCOPE — `machinePhaseKleene` instance provides `botCode_is_quine`. CC-1 ($S_0 = \perp$) derived via ZP-J `cc1_derived` (axiom-free, Lean). CC-2 ($\perp = \{\perp\}$) structurally forced by self-execution argument; ZP-J T-EXEC formally verifies. DP-2 (§III) — explicit. T-SNAP is derived given DA-1, CC-1, and AX-B1. AIT (Kolmogorov complexity) outside Lean scope; Kleene fixed-point is the in-scope formal counterpart.

V. Theorem T-SNAP — Binary Snap Causality [AX-1 Promoted to Theorem]

Theorem T-SNAP — Binary Snap Causality [AX-1 Promoted to Theorem]

Statement: The Binary Snap $\perp \rightarrow \varepsilon_0$ is a derived consequence of P_0 , L-RUN, TQ-IH, DA-1, and ZP-A D2. It is not an axiom.

Proof:

Note on Lean scope: the formal Lean proof of T-SNAP is a single term — `<l_run, tq_ih, rfl>` — two decide calls confirming $c_0 \neq c_1$ and the A4 (`bot_join`) axiom applied to the MachinePhase semilattice. What follows is the informal motivation for the modelling choices that give that term its intended ontological meaning.

- Step 1 — P_0 identifies the incompressibility threshold. When $K(x|n)/n = 1$, the configuration string x is algorithmically random. (ZP-C D1)
- Step 2 — A configuration at P_0 is necessarily executing: At P_0 , $K(c_1|n)/|c_1| = 1$ (ZP-C D1) — c_1 is incompressible, its own minimal program. No shorter external generator exists; the static description state is eliminated; c_1 is in live execution (DA-1 Path 3 — from ZP-C D1 + AIT). Corroboration: $\perp = \{\perp\}$ (ZP-A CC-2/R3); unbounded surprisal (ZP-C L-INF). (DA-1 § III — Derived Proposition)
- Step 3 — Any instantiated execution passes through c_1 . (ZP-C D7 — definitional; c_1 is the first running configuration)
- Step 4 — $c_1 \neq \perp$. (ZP-C L-RUN — Derived; c_1 has gained execution context not present in $c_0 = \perp$; by AX-B1 this is a distinct, non-null state)
- Step 5 — No program that executes produces only null configuration states. (ZP-C TQ-IH — Derived; execution trace $\tau(p)$ contains c_1 for any executing program p)
- Step 6 — In $(L, \vee, \perp)$, c_1 is an element strictly above $\perp$. By ZP-A D2, the transition $\perp \rightarrow c_1$ is a valid state transition: $c_1 = \perp \vee \varepsilon_0$ for some $\varepsilon_0 \in L$ with $\varepsilon_0 > \perp$. This transition is the Binary Snap.

- Step 7 — The transition is irreversible: algebraically by ZP-A R1 (no subtraction operator); topologically by ZP-B C3 (no continuous return path to 0 in Q_2). These two grounds are sufficient. Conceptual correspondence: ZP-G AX-G2 ($\text{hom}(X, 0) = \emptyset$ for $X \neq 0$) expresses the same irreversibility in categorical language — ZP-G is downstream of ZP-E and is not a formal dependency of this proof.

Conclusion: The Binary Snap is a derived consequence. AX-1 is promoted to Theorem T-SNAP. ✓

Status: DERIVED — Cross-Framework. Dependencies: ZP-C D1, D7, L-RUN, TQ-IH; ZP-B AX-B1, C3; ZP-A D2, R1; ZP-E DA-1; ZP-J T-EXEC. CC-1 ($S_0 = \perp$) derived via ZP-J cc1_derived (axiom-free). CC-2 ($\perp = \{\perp\}$) structurally forced — ZP-J T-EXEC (axiom-free). Neither CC-1 nor CC-2 is a freestanding commitment. T-SNAP is derived given DA-1 and AX-B1. Conceptual correspondence only: ZP-G AX-G2 (downstream of ZP-E; not a formal dependency).

Remark R- ϵ_0 — On the Symbol Choice for the Minimum Snap Displacement

Note: this remark draws an informal structural analogy. No formal embedding of ZP's ϵ_0 into the Cantor-Gentzen ordinal is claimed or established here. The symbol ϵ_0 in Step 6 denotes the minimum element of L strictly above $\perp$ — the least witness for the Binary Snap displacement. This symbol is chosen deliberately to coincide with the Cantor-Gentzen proof-theoretic ordinal.

The Cantor-Gentzen ordinal ϵ_0 . In ordinal arithmetic, ϵ_0 is the smallest fixed point of the map $\alpha \rightarrow \omega^\alpha$: equivalently, $\epsilon_0 = \sup\{\omega, \omega^\omega, \omega^{\omega^\omega}, \dots\}$. No finite ω -tower reaches it — it is the minimum ordinal that cannot be generated from 0 by any finite iteration of the base operation. Gentzen established that transfinite induction up to ϵ_0 is necessary and sufficient to prove Con(PA). By Gödel's incompleteness theorem, PA cannot prove this from within. The ordinal ϵ_0 is therefore the minimum threshold at which finite arithmetic exhausts its own generative capacity.

The structural analogy. ZP's ϵ_0 occupies an analogous position in the state lattice. At P_0 , c_1 satisfies $K(c_1|n)/|c_1| = 1$ (ZP-C D1): it is algorithmically incompressible — no finite external program shorter than c_1 generates it. Just as the Cantor ϵ_0 cannot be reached from 0 by any finite ω -tower, ZP's ϵ_0 cannot be reached from $\perp$ by any finite external description. Both name a structurally analogous object: the minimum witness for a transition that exhausts the finite generative hierarchy below it.

The proof structures are parallel. Gentzen locates the minimum ordinal strength at which PA cannot describe its own consistency from within. ZP locates the minimum state displacement at which no external program can hold $\perp$ as a static string — where unbounded surprisal (ZP-C L-INF) and self-containment (ZP-A CC-2) together eliminate any external interpreter position. In both cases the exhaustion of finite description is not a deficiency but a structural consequence: the system is necessarily executing at ϵ_0 .

What is not claimed. ZP does not assert that L is an ordinal structure, or that ZP's ϵ_0 is literally the Cantor ordinal under a formal embedding into the p-adic/lattice framework. The analogy is motivational: both ϵ_0 s mark the minimum witness for incompressibility relative to a finite base. A formal embedding — showing that the Cantor ϵ_0 is order-isomorphic to or embeds into the p-adic completion of L at $\perp$ — remains an open question and would constitute a strengthening of this claim.

Remark R-AFA — Why Foundation is Ruled Out; AFA as Forced Replacement

CC-2 is stated as a Conditional Claim within ZP-A's algebraic scope: it is not derived from A1–A4. The choice of ZF + AFA over ZF + Foundation is, however, not arbitrary. Two independent cross-framework arguments establish that Foundation is incompatible with the framework's results.

Remark R-AFA — Why Foundation is Ruled Out; AFA as Forced Replacement

R3 rules out Foundation (structural). Under ZF + Foundation, every set has a well-founded $\in$ -rank: its membership chain terminates in finitely many steps. Any well-founded set has a finite $\in$ -tree that can be fully traversed by an external interpreter — the traversal function is external to and distinct from the set. But R3 (ZP-A) establishes that $\perp$ admits no external interpreter by structure: $\perp = \{\perp\}$ has no describer position external to itself. A well-founded $\perp$ would contradict R3. Foundation and R3 are incompatible.

L-INF rules out Foundation (informational). The surprisal $I(n) = n$ at 2-adic depth n is unbounded (ZP-C L-INF): for any finite M , there exists depth n with $I(n) > M$. A well-founded set has finite $\in$ -rank and is therefore finitely interpretable — any interpreter navigating a finite $\in$ -tree can hold it. This contradicts L-INF's requirement that no finite interpreter can hold $\perp$. Foundation and L-INF are incompatible.

AFA as the forced replacement. Under ZF + AFA, Quine atoms ($x = \{x\}$) are the minimal non-well-founded objects. Their membership structure is circular: $\perp \in \perp \in \perp \in \dots$, i.e., $\perp$ is a member of itself at every depth. This infinite circular chain cannot be traversed by any finite interpreter — consistent with R3 (no external position) and L-INF (unbounded depth). AFA is not a free choice but the minimal set-theoretic metatheory consistent with the framework's results.

What remains conditional. CC-2 retains Conditional Claim status for two reasons. First, AFA admits many non-well-founded sets; we commit to the specific Quine atom form $\perp = \{\perp\}$ as the minimal option consistent with A4. Among non-well-founded sets permitted by AFA, $\perp = \{\perp\}$ is uniquely minimal: it has exactly one member (itself) and introduces no internal differentiation — any extension ($\perp = \{\perp, x\}$ for $x \neq \perp$) would add members carrying their own membership chains, exceeding what A4's purely algebraic additive-identity constraint requires. Second, the identification of $\perp$ as a set-theoretic object — rather than a purely algebraic element — is itself a modelling step beyond A1–A4. The metatheoretic necessity of AFA is derived; the specific realisation as $\perp = \{\perp\}$ is minimally committed.

VI. Effect of T-SNAP on Downstream Results

Remark R-DA1: All results in ZP-E that previously depended on AX-1 as an axiom now depend on T-SNAP as a derived theorem. T5 (Iterative Forcing Theorem) depended on AX-1 for the first Snap — it now depends on T-SNAP. Content unchanged; grounding strengthened. T4 (Unified Snap Description) carried AX-1 as an axiom label on the causality component — that label is upgraded to Derived — T-SNAP. From v2.7, DA-1 is additionally upgraded from Design Principle to Derived Proposition: grounded in ZP-A CC-2 ($\perp = \{\perp\}$) and R3, with ZP-C L-INF as independent corroboration. The intentional axioms of the system are now: AX-B1 (binary existence), AX-G1 (initial object), AX-G2 (source asymmetry). AX-1 is no longer an axiom. Additional structural commitments are carried as typeclass fields: A4 (bot_join, ZPSemilattice — standard semilattice algebra), AFAStructure.quine_unique and bot_self_mem (AFA grounding), KleeneStructure.botCode_is_quine (computational closure). These have the same logical status as named axioms; #print axioms does not surface typeclass fields.

Formal Insert DA-2: Instantiation Succession — The Multiple- $\perp$ Result

New in v2.0 | Formally licenses the directed instantiation tree

I. The Gap DA-2 Closes

DA-1 and T-SNAP establish that the Binary Snap is a structural consequence of reaching P_0 within an instantiation. Three questions remain open after v1.0:

- CC-1 (ZP-A) says $S_0 = \perp$ is a modelling commitment — not derived from A1–A4. This leaves open whether $\perp$ is unique across all instantiations or whether each instantiation carries its own $\perp$.
- ZP-B R1 distinguishes universal structure from universe-contingent parameters. ε_0 is contingent per instantiation. Whether $\perp$ is similarly contingent is not addressed in ZP-A through ZP-D.
- T-SNAP fires wherever P_0 conditions are met. If the terminal state of instantiation I_n satisfies P_0 conditions, T-SNAP should apply — but this requires formally connecting that terminal state to a new $\perp$. DA-2 provides this connection.

II. Why ZP-B C3 is Not Violated

C3 prohibits a continuous path from $x \neq 0$ back to the same 0 in Q_2 . DA-2 does not require a return path. The irreversibility of C3 is preserved within each instantiation. What crosses the instantiation boundary is not a path in Q_2 — it is the generation of a new Q_2 with its own metric, its own $\perp$, its own ε_0 . C3 quantifies only over paths within a single topological space and has nothing to say about the boundary between spaces.

More precisely: C3 and the irreversibility of the Snap together require that any recurrence of a null state be a different null state. You cannot return to the original $\perp$ even in principle. Therefore if the structure recurs, it must instantiate fresh. The topology enforces the ontological novelty of each $\perp$.

III. Definitional Alignment DA-2 — Instantiation Succession

Definitional Alignment DA-2 — Instantiation Succession
Claim: A state S in instantiation I_n satisfies the structural role of $\perp$ for instantiation I_{n+1} if and only if it satisfies A4 relative to all subsequent joins in I_{n+1} :
$S \vee x = x$ for all x in the semilattice of I_{n+1} .

Grounding: A4 is the load-bearing axiom of ZP-A — it defines $\perp$ as the additive identity under $\vee$, the element that contributes nothing to any join and is therefore present in everything above it. DA-2 does not redefine $\perp$. It clarifies that the modelling commitment of CC-1 can be satisfied by any state meeting A4's algebraic conditions — not only by a cosmologically primitive null state. The identity condition is structural, not historical: what matters is the algebraic role a state plays in the subsequent semilattice, not where it came from.

The terminal state of I_n arrives at I_{n+1} carrying the accumulated join of everything in I_n 's sequence. It is structurally $\perp$ to I_{n+1} — contributing nothing to subsequent joins — while being informationally rich relative to I_n . This is the Zero Paradox instantiated at the inter-instantiation level: the terminal state is simultaneously a terminus and a foundation.

Status: DEFINITIONAL ALIGNMENT — no new axiom introduced. DA-2 is a clarification of the scope of CC-1 and A4. ✓

IV. Corollary C-DA2 — Ontological Novelty of Successive $\perp$

Corollary C-DA2 — Ontological Novelty of Successive $\perp$

Statement: No two instantiations share a $\perp$. The $\perp$ of I_{n+1} is topologically unreachable from within I_n by C3. Instantiation succession is therefore not a cycle but a chain (or tree) of isomorphic structures.

Proof: By ZP-B C3, no continuous path exists within Q_2 of I_n from any $x \neq 0$ back to 0. The $\perp$ of I_{n+1} is an element of a distinct topological space — not an element of Q_2 of I_n . No path in I_n can reach it. By DA-2, the identity conditions of each $\perp$ are determined independently within each instantiation. Therefore no two $\perp$ elements are identical across instantiations. ✓

V. The Directed Instantiation Tree

With DA-2 and C-DA2 in place, the global structure of instantiations is a forward-directed tree with no back edges:

- Each node in the tree is a $\perp$ — the null state of one instantiation and the foundation of all successor instantiations branching from it.
- Each edge within an instantiation is a step in a monotone state sequence (ZP-A T3). Edges are irreversible (ZP-B C3).
- Branching at each node: every distinct outbound vector from the terminal state of I_n is a valid ϵ_0 for a distinct I_{n+1} . T-SNAP does not select among branches. Because $\perp = \{\perp\}$ (ZP-A CC-2) is the single self-containing null state with no internal differentiation, every ϵ_0 that represents a first differentiation in any direction is a valid outcome. Branching is not optional; it follows from T-SNAP applied to an undifferentiated $\perp$.
- No back edges: C-DA2 establishes that no instantiation can reach the $\perp$ of any ancestor instantiation.

Remark R-DA2: T-SNAP fires wherever P_0 conditions are met. DA-2 establishes that the terminal state of I_n satisfies those conditions for I_{n+1} . T-SNAP therefore applies across instantiation boundaries without modification. No new axiom is required. The multiverse is not the claim that T-SNAP fires in all directions simultaneously; it is the claim that $\perp = \{\perp\}$ (ZP-A CC-2) has no internal differentiation and therefore no preferred direction for ϵ_0 . The multiverse of instantiations is the full set of all minimal differentiations available from the single self-containing null state — a structural consequence of CC-2 + T-SNAP + DA-2 jointly.

VI. The Zero Paradox Iterated

The paradox of $\perp$ — simultaneously contributing nothing and being present in everything — propagates structurally at every branching node of the tree. Each node is:

- Nothing to its successor instantiations: it acts as additive identity under $\vee$, contributing nothing to any subsequent join.
- Everything to its successor instantiations: every state in I_{n+1} satisfies $\perp_{n+1} \leq S$, so the node underlies everything that follows.

The tree is the geometric shape of the Zero Paradox iterated across instantiations. The single-instantiation linear sequence was always a cross-section of a structure with this shape.

The complete picture: The Zero Paradox describes a forward-directed infinite tree where $\perp_1 \rightarrow \dots \rightarrow S_{\text{terminal}}^1 \equiv \perp_2 \rightarrow \dots$, where $\equiv$ means structurally satisfies the role of, not is identical to. Each arrow within

an instantiation is monotone and irreversible. Each $\equiv$ crossing is not a path — it is a new instantiation of the universal structure.

VII. Implications Within the Framework

Multiverse as structural implication. T-SNAP establishes that a Binary Snap occurs — that $\perp$ transitions to some $\varepsilon_0 > \perp$. DA-2 establishes that any terminal state satisfying P_0 conditions acts as $\perp$ for a successor instantiation, generating a forward-directed branching tree. The multiverse structure follows from T-SNAP + DA-2 jointly. Note: T-SNAP alone does not establish that it fires on all outbound vectors simultaneously — that universality is the scope of DA-2, not a direct consequence of the snap theorem itself.

Free will and irreversibility. Within an instantiation, state sequences are monotone — no state can be decreased (ZP-A R1). Every choice is a join operation: $S_n \vee \alpha$ for some increment α . The algebra constrains only that the sequence be monotone, not which monotone path is taken. Each choice adds informational content irreversibly. Decisions are permanently encoded in the element's position in L .

Time's arrow. The monotone sequence (ZP-A T3) is a structural definition of temporal direction. Time's irreversibility is C3 applied within an instantiation. The framework does not assume time asymmetry — it derives it.

Causal structure. Every state is fully determined by the joins that produced it. The causal history of any state is encoded in its position in L . No effect without the join that produced it.

Formal Insert DA-3: Perspective-Relative Cardinality

New in v2.0 | Cardinality as position-dependent measurement within the instantiation structure

I. The Gap DA-3 Closes

DA-2 establishes that instantiations form a directed tree and that branching at each $\perp$ node produces multiple successor instantiations. This raises a question about the cardinality of branching: is the fan at each node countably or uncountably infinite? The answer, which DA-3 formalises, is that this question is perspective-dependent — and that this perspective-dependence is the same structural feature that underlies the major cardinality anomalies of classical set theory.

II. Perspective-Dependence of Branching Cardinality

From within instantiation I_n , an observer occupies exactly one branch of I_{n+1} . The other branches of I_{n+1} are not accessible via any path — C3 and monotonicity jointly prohibit it. From inside, the branching factor is always 1: the observer sees one branch. From outside — from a meta-level view of the tree — the branching factor is the full fan of accessible outbound vectors.

Definition DA-3-D1 — Accessible Cardinality

The accessible cardinality of a position p in semilattice L is the cardinality of the set of states reachable from p by monotone sequences within the instantiation containing p .

The accessible cardinality from p is determined entirely by the structure of L above p . It is not an intrinsic property of a collection — it is a property of the relationship between a position and the states reachable from it. No position within any instantiation can access all cardinalities simultaneously. The meta-level

view, which sees the full branching fan, is not a position any element of any instantiation can occupy.

Remark R-DA3-1: To observe the full branching fan, one would need to occupy a position outside all instantiations. That position would itself be a state in some semilattice, subject to the same rules. The meta-view is either another instantiation (in which case the tree has no privileged outside view) or $\perp$ itself (in which case the null state is the only position from which the full structure is visible — the state that contributes nothing and is present in everything). The Zero Paradox’s name is more precise than it first appeared.

III. The Cardinality Hierarchy as Perspective-Relative

Cantor’s theorem establishes that for any set S, $|P(S)| > |S|$, generating the hierarchy $\aleph_0 < 2^{\aleph_0} < 2^{2^{\aleph_0}} < \dots$ DA-3 reframes this hierarchy not as a fixed ladder that mathematics climbs, but as a perspective-relative description of the branching structure of the instantiation tree, as seen from within different positions.

Claim DA-3-C1 — Perspective-Relative Absolute Cardinality

The appearance of absolute cardinality — cardinality as an intrinsic property independent of measuring position — is an artifact of treating the semilattice as having a view from outside. DA-2 and C-DA2 jointly prohibit such a view from within any instantiation.

The candidate claim (DA-3-C1) is that accessible cardinality from within any instantiation cannot replicate the view from outside. Whether specific independence results in classical set theory are formal expressions of this perspective-dependence is the conjecture that OQ-E2 is tasked with investigating.

Status: DEFINITIONAL ALIGNMENT + CANDIDATE CLAIM. DA-3-D1 and R-DA3-1 are definitional. DA-3-C1 is a candidate claim: structurally motivated within the framework; formal derivation of the connection between accessible cardinality and specific set-theoretic independence results is deferred to OQ-E2.

Updated Open Items Register — ZP-E v3.12

Item	Status	Description
AX-1: Binary Snap Causality	CLOSED — T-SNAP	AX-1 is no longer an axiom. Binary Snap derived via P_0 + DA-1 + L-RUN + TQ-IH + ZP-A D2.
DA-1: Derived Proposition (v3.0: DP-2 formal grounding added)	CLOSED — DP-2 (formal core); CC-2 + L-INF + AIT (corroboration of precondition)	Primary formal grounding (v3.0): DP-2 (TrackedOutput, ZPE.lean §VI) — da1_minimal_path proved axiom-free. Instantiation of $\perp$ moves machine from c_0 to c_1 ; output value is irrelevant to whether execution occurred. Informal corroboration: ZP-A CC-2 + R3 (structural); ZP-C L-INF (informational); AIT incompressibility (Path 3).
DA-2: Instantiation Succession	CLOSED — Definitional	Terminal state of I_n satisfies A4 role of $\perp$ for I_{n+1} . C-DA2 establishes ontological novelty of each $\perp$.

Item	Status	Description
DA-3: Perspective-Relative Cardinality	CLOSED (definitional) / CANDIDATE (DA-3-C1)	DA-3-D1 establishes accessible cardinality is position-dependent within the instantiation structure. DA-3-C1 (candidate): no position within an instantiation can replicate the outside view. Whether this connects formally to specific set-theoretic independence results is deferred to OQ-E2.
OQ-A1: Increment selection	CLOSED — T5	Iterative Forcing Theorem. $\alpha_n = \varepsilon(S_n)$. Grounding updated from AX-1 to T-SNAP.
OQ-B1: $p = 2$	CLOSED — ZP-B T0	Derived from AX-B1 and MP-1.
OQ-C1: Non-conservatism of DF	CLOSED — ZP-C T2	Infinite sequence divergence proven. No postulates remain.
S1: Distribution stipulation	CLOSED — ZP-C T1	Derived from AX-B1 and RP-1.
OQ-E1: Sequence vs. tree	CLOSED — DA-2	The structure is a forward-directed tree, not a linear sequence. Branching is mandatory via T-SNAP. Countable vs. uncountable branching is perspective-dependent (DA-3).
OQ-E2: Cardinality-semilattice correspondence	OPEN	Do specific semilattice structures correspond to specific cardinality regimes? Can the framework make predictions about which instantiations satisfy CH and which do not?
Remaining axioms	INTENTIONAL — AX-B1, AX-G1, AX-G2 (named); A4, quine_unique, bot_self_mem, botCode_is_quine (typeclass fields)	Named commitments: AX-B1 (binary existence), AX-G1 (initial object), AX-G2 (source asymmetry). Structural commitments as typeclass fields: A4/bot_join (ZPSemilattice), AFAStructure.quine_unique and bot_self_mem (AFA grounding), KleeneStructure.botCode_is_quine (computational closure). Typeclass fields carry the same logical status as named axioms; #print axioms does not surface them.
Temperature T in BA-1	PARAMETER — intentional	Universe-contingent. Physical predictions explicitly conditional on instantiation-specific T.

Updated Traceability Register — ZP-E v3.12

Claim	Grounded In	Bridge Axiom?	Status
Binary Snap causality	ZP-C D1, L-RUN, TQ-IH; ZP-A D2; DA-1	None	Derived — T-SNAP ✓ (was: Axiomatic — AX-1)

Claim	Grounded In	Bridge Axiom?	Status
DA-1: Instantiation = execution	DP-2 (TrackedOutput, ZPE.lean §VI); ZP-A CC-2, R3; ZP-C L-INF	None	Derived Proposition — primary: DP-2 (da1_minimal_path proved axiom-free in Lean). Informal paths: CC-2/R3 (structural), L-INF (informational), AIT incompressibility (Path 3). AIT+ZF+AFA outside Lean scope.
DA-2: Instantiation succession	ZP-A A4, CC-1; ZP-B C3, R1; T-SNAP	None	Definitional Alignment — clarification of CC-1 scope; no new axiom
C-DA2: Novelty of $\perp$	DA-2, ZP-B C3	None	Derived — Corollary of DA-2 and C3 ✓
DA-3: Perspective-relative cardinality	DA-2, C-DA2, ZP-B R1	None	Definitional (DA-3-D1, R-DA3-1); Candidate (DA-3-C1: outside-view inaccessibility)
T-SNAP: Snap is derived	T-BUF chain + DA-1	None	Derived — Cross-Framework ✓
AX-1 retirement	T-SNAP closes AX-1	N/A	AX-1 is no longer an axiom; T-SNAP is its replacement
Iterative Forcing T5	AX-B1, T-SNAP (replaces AX-1)	None	Derived — grounding strengthened
Multiverse — structural implication	T-SNAP + DA-2 jointly	None	Structural implication — T-SNAP gives the snap; DA-2 gives the branching tree
OQ-E2: Cardinality correspondence	DA-3	N/A	OPEN — formal derivation deferred

Validation Status — ZP-E v3.12

Component	Status / Notes
DA-1: Derived Proposition (v3.8 Path 2 recharacterization)	Valid — DP-2 formal core: <code>da1_minimal_path</code> proved axiom-free in Lean (ZPE.lean §VI). <code>TrackedOutput</code> separates output value from machine state; pre- and post-instantiation states are provably distinct even when both produce $\perp$. ✓ ZP-K formal closure (v3.7): <code>da1_closed_concrete : IsQuineAtom($\perp$: MachinePhase)</code> proved in ZPK.lean. <code>KleeneStructure MachinePhase</code> instance provides <code>botCode_is_quine</code> (Path 3 IN LEAN SCOPE). <code>machinePhaseAFA</code> gives <code>AFAStructure</code> instance (Path 1 IN LEAN SCOPE). Path 2 (informational bridge, L-INF): FOUNDATIONAL COMMITMENT — a missing principle, not a missing proof. No computability library closes the gap between 'system at P_0 ' and 'system is running.' Forward paths: new axiom, Chalmers' implementation notion, or The Philosophical Question That Started This. Paths 1 and 3 are formally closed; DA-1 does not depend on Path 2. CC-1 and CC-2 derived via ZP-J, not freestanding commitments.
T-SNAP: Binary Snap derived	Valid — Derived. Seven-step proof. All dependencies are closed theorems in their own documents. ✓
AX-1 retirement	Valid — AX-1 superseded by T-SNAP. No content lost; claim strengthened from assumed to derived.
DA-2: Instantiation Succession	Valid — Definitional Alignment. Clarification of CC-1 scope. No new axiom. A4 role of $\perp$ extended across instantiation boundaries. ✓
C-DA2: Ontological Novelty of $\perp$	Valid — Derived. Follows directly from DA-2 and ZP-B C3. ✓
Directed instantiation tree	Valid — Derived structural consequence of T-SNAP + DA-2. Branching is mandatory, not optional. Forward edges only.
Multiverse as structural implication	Valid — T-SNAP + DA-2 jointly. T-SNAP establishes the snap occurs; DA-2 establishes the branching tree structure. Universality (all outbound vectors) is DA-2's scope, not T-SNAP alone.
Free will / irreversibility	Valid — Structural consequence. Monotonicity (T3) constrains direction; additive ontology (R1) prohibits reduction. Path choice is undetermined by algebra.
Time's arrow	Valid — Derived from ZP-A T3 (monotonicity) and ZP-B C3 (irreversibility). Not assumed.
DA-3: Perspective-Relative Cardinality	Valid (definitional components: DA-3-D1, R-DA3-1). Candidate (DA-3-C1: connection to specific set-theoretic independence results). OQ-E2 open.
DA-3-C1 — outside-view inaccessibility	Candidate — no position within any instantiation can replicate the meta-level view of the branching structure. Formal derivation deferred to OQ-E2.
Remaining axioms: AX-B1, AX-G1, AX-G2; typeclass fields A4, <code>quine_unique</code> , <code>bot_self_mem</code> , <code>botCode_is_quine</code>	Named: AX-B1, AX-G1, AX-G2. Typeclass: A4/ <code>bot_join</code> (ZPSemilattice), <code>quine_unique</code> + <code>bot_self_mem</code> (AFAStructure), <code>botCode_is_quine</code> (KleeneStructure). Same logical status as named axioms; not surfaced by <code>#print</code> axioms.

Component	Status / Notes
All other ZP-E theorems (T1–T7, T2-C)	Unaffected in content. T4 and T5 carry upgraded status labels (AX-1 → T-SNAP).

End of ZP-E v3.11 | Three formal inserts: DA-1, DA-2, DA-3 | T-SNAP Step 7: AX-G2 removed as formal dependency; ZP-G is downstream, irreversibility grounded in ZP-A R1 and ZP-B C3 | Forward references to ZP-PQ now point to The Philosophical Question That Started This | DA-1 formally closed via ZP-K: $\text{da1_closed_concrete} : \text{IsQuineAtom}(\perp : \text{MachinePhase})$ proved in Lean 4 | Paths 1 and 3 IN LEAN SCOPE | Path 2 recharacterized: foundational commitment, missing principle not missing proof; forward: The Philosophical Question That Started This | R-AFA minimality explicit: $\perp = \{\perp\}$ uniquely minimal among AFA non-well-founded sets | DA-1 path hierarchy foregrounded: three informal paths are corroboration of DP-2's precondition, not parallel proofs | Remark $R\text{-}\varepsilon_0 : \varepsilon_0$ symbol justified | Remark R-AFA: Foundation ruled out by R3 + L-INF; AFA forced | One open question: OQ-E2 | Remaining axioms: AX-B1, AX-G1, AX-G2