

THE ZERO PARADOX

ZP-O: The Fixed-Point Fork

Version 1.0 | June 2026

v1.0: Initial release. Synthesis layer. The abstract fork schema is proved sorry-free and choice-free in Lean 4 (ZPO.lean); the number-system instance is discharged through Ostrowski's theorem (ZPO_Ostrowski.lean). Generalizes the ZFC+Foundation / ZFC+AFA orthogonal-contact-point claim.

A self-referential operator over a complete lattice has a least fixed point and a greatest fixed point (Knaster–Tarski). ZP-O records a single elementary fact and develops its consequences across the framework: these two fixed points coincide — the fork collapses to one point — exactly when the operator has a unique fixed point. Read intuitively, the least fixed point is the inductive (well-founded) closure and the greatest is the coinductive (non-well-founded) closure; the collapse point is where the two closures meet. In the Zero Paradox that point is the diagonal fixed point $\perp$, the self-containing bottom.

This layer is a synthesis layer, not a foundational one: it adds no axiom or primitive. It names a pattern that several existing layers already realise, and it is organised in three tiers held to three different standards of proof. Tier 1 (Section I) is the abstract schema, proved in Lean over a complete lattice. Tier 2 (Section II) is the instance catalogue: each framework forks at its own contact point, and the instances are theorems in their home layers. Tier 3 (Section III) is the unification — that the contact points are faces of one object — which is a fenced conjecture, never a formal claim. The fences in Section III are load-bearing: keeping the tiers separate is what makes the formal content honest.

Section I: The Fork Schema (Tier 1)

Over a complete lattice, a monotone self-map f has a least fixed point ($\text{lfp } f$) and a greatest fixed point ($\text{gfp } f$), and $\text{lfp } f \leq \text{gfp } f$ always. The width between them is the fork. The schema theorem states that the fork collapses to a single point precisely when f has a unique fixed point — and in that case the common value is that fixed point. The proofs rest only on Mathlib's order-theoretic fixed-point API (`OrderHom.lfp` / `OrderHom.gfp`); they make no claim about the categorical inductive/coinductive reading, which is offered as analogy only.

Definition: the fork (ZPO.lean)

For a complete lattice α and a monotone self-map $f : \alpha \rightarrow \alpha$:

$\text{lfp } f$ = the least fixed point of f (Mathlib `OrderHom.lfp`)

$\text{gfp } f$ = the greatest fixed point of f (Mathlib `OrderHom.gfp`)

The fork is the ordered pair $(\text{lfp } f, \text{gfp } f)$, always satisfying $\text{lfp } f \leq \text{gfp } f$.

Collapse = $\text{lfp } f = \text{gfp } f$: the two ends meet at a single contact point.

Proposition: fork_le (ZPO.lean)

$\text{lfp } f \leq \text{gfp } f$

The inductive closure never exceeds the coinductive closure: the fork has non-negative width. (Mathlib OrderHom.lfp_le_gfp.)

Lean purity: [propext, Quot.sound] — choice-free. ✓

Lemma: collapse_of_unique (ZPO.lean)

If x is the unique fixed point of f , then $\text{lfp } f = x$ and $\text{gfp } f = x$.

Both ends of the fork land on the sole fixed point.

Proof: $\text{lfp } f$ and $\text{gfp } f$ are fixed points (map_lfp , map_gfp); uniqueness sends each to x .

Lean purity: [propext, Quot.sound] — choice-free. ✓

Lemma: unique_of_collapse (ZPO.lean)

If $\text{lfp } f = \text{gfp } f$, then every fixed point of f equals that common value.

Every fixed point lies between $\text{lfp } f$ and $\text{gfp } f$, so collapse forces uniqueness.

Proof: lfp_le_fixed and le_gfp bracket any fixed point y ; antisymmetry closes it.

Lean purity: [propext, Quot.sound] — choice-free. ✓

Theorem: fork_collapse_iff (ZPO.lean) — the schema spine

$\text{lfp } f = \text{gfp } f \leftrightarrow \exists! x, f x = x$

The fork collapses to a single contact point if and only if f has a unique fixed point.

This is the abstract form of "the framework forks, and the diagonal fixed point is where the two closures meet."

Proof: $\text{collapse_of_unique}$ and $\text{unique_of_collapse}$, both directions.

Lean purity: [propext, Quot.sound] — choice-free. ✓

The fork spine is choice-free. All four results depend only on [propext, Quot.sound] — propositional extensionality and quotient soundness, the benign kernel axioms used throughout Mathlib. None depends on the Axiom of Choice. This is consistent with the choice-free core (see AxiomProfile.lean): the conceptual skeleton needs no choice; choice enters only in the analytic realisations (Section II).

Section II: The Instance Catalogue (Tier 2)

The same fork appears across frameworks. ZP-O does not re-prove each instance — each is a theorem in its home layer — it records the catalogue and pins the number-system instance to a checkable Lean witness. The original ZFC+Foundation / ZFC+AFA orthogonal-contact-point claim is the set-theory entry: in the standard category-theoretic characterisation (Aczel; Lambek), Foundation is the initial algebra (μ) and AFA the final coalgebra (ν) of the bounded powerset functor, and the Quine atom $\perp = \{\perp\}$ is exactly an element present in ν but not in μ .

Domain	The two closures fork into	Contact point / home layer
Set theory	Foundation (μ , well-founded) vs AFA (ν , non-well-founded)	Quine atom $\perp = \{\perp\}$; ZP-J / ZP-K
Number systems	Archimedean $\mathbb{R}$ vs non-Archimedean $\mathbb{Q}_p$	0 (snap fails in $\mathbb{R}$, holds in $\mathbb{Q}_2$); ZP-B / ZP-F
Computation	total (well-founded) vs partial (non-terminating)	the Kleene quine / self-application; ZP-K
Category theory	initial algebra (W-types) vs final coalgebra (M-types)	QPF.Fix vs QPF.Cofix (the parent schema)

The number-system instance is the one with a genuine classification theorem behind it. Ostrowski’s theorem says the nontrivial absolute values on $\mathbb{Q}$ are exactly the real (Archimedean) one and the p-adic (non-Archimedean) ones — the complete, pairwise-inequivalent list. The contact point is 0: in the Archimedean completion $\mathbb{R}$ it is approached but never reached (density — the snap fails), while in the 2-adic completion $\mathbb{Q}_2$ it is the unique fixed point of $x \mapsto 2x$ and its valuation diverges ($v_2(0) = \infty$ — the snap holds).

Theorem: completions_exhaustive (ZPO_Ostrowski.lean)
Every nontrivial absolute value on $\mathbb{Q}$ is equivalent to the real absolute value, or to a p-adic absolute value for a unique prime p.
The two kinds of completion are the complete list. (Ostrowski’s theorem, Mathlib Rat.AbsoluteValue.equiv_real_or_padic.)
Lean purity: [propxt, Classical.choice, Quot.sound] — choice-carrying. ✓

Theorem: real_not_equiv_padic (ZPO_Ostrowski.lean)
The real absolute value is inequivalent to every p-adic absolute value.
The two kinds of completion are genuinely distinct — never the same metric. This is the orthogonality leg of the number-system fork. (Mathlib Rat.AbsoluteValue.not_real_isEquiv_padic.)
Lean purity: [propxt, Classical.choice, Quot.sound] — choice-carrying. ✓

Remark: choice-free spine, choice-carrying realisation
The fork spine (Section I) is choice-free [propxt, Quot.sound]. The Ostrowski instance carries Classical.choice, inherited from Mathlib’s classical analysis and number theory. This split is the framework’s standing pattern: the conceptual core is choice-free, while the analytic realisations are choice-carrying.

Remark: choice-free spine, choice-carrying realisation

The number-system fork is theorem-backed on its own terms — Ostrowski is a genuine classification theorem, stronger backing than the metatheoretic Foundation/AFA orthogonality. It is NOT claimed to be an instance of the μ/ν (least-vs-greatest fixed point) schema: Ostrowski concerns absolute values, not fixed points of a functor. The thread to the diagonal fixed point runs through the contact point 0 (the unique fixed point of $x \mapsto 2x$ in $\mathbb{Q}_2$), not through the schema theorem.

Section III: The Unification and Its Fences (Tier 3)

The tempting claim is that the contact points of all the instances — the Quine atom, the 2-adic 0, the Kleene quine, the categorical initial object — are faces of one object, the diagonal fixed point. ZP-O states this as a conjecture and fences it precisely. The fences are not hedging; they mark a genuine boundary, and per-instance forks remain full theorems on either side of them.

Hard fence (permanent): cross-instance identity is not a theorem

The claim "the contact points are one object" is not a theorem and cannot become one. Identity requires a shared type: the 2-adic 0, the Quine atom (a set), the Kleene quine (a code), and a categorical initial object (an object up to isomorphism) are terms of different types in different categories. The proposition " $x = y$ " across them is not false — it is not well-formed. So this is a type boundary, not a missing proof. What is claimed formally is only that each framework forks at its own contact point.

Soft fence (conjecture): not every fork is a μ/ν instance

That every domain fork is an instance of the least-vs-greatest fixed point schema is an open generalisation, not established here. The schema is a theorem in this layer and for canonical functors (W-types vs M-types). The number-system fork (Ostrowski) is the clear case that is theorem-backed yet NOT visibly a μ/ν instance — it is the live caution against overstating the generalisation.

Per-instance forks are theorems and are not hedged. The fences scope only the cross-instance identity (Tier 3) and the universal schema claim. Each individual fork — Foundation/AFA, $\mathbb{R}/\mathbb{Q}_2$, total/partial — stands on its own proof.

Theorem Summary

Result	File / Section	Axioms
fork_le	ZPO.lean §I	choice-free
collapse_of_unique	ZPO.lean §I	choice-free
unique_of_collapse	ZPO.lean §I	choice-free
fork_collapse_iff	ZPO.lean §I	choice-free
completions_exhaustive	ZPO_Ostrowski.lean §II	Classical.choice

Result	File / Section	Axioms
real_not_equiv_padic	ZPO_Ostrowski.lean §II	Classical.choice

Axiom Purity
Fork spine (ZPO.lean): [propext, Quot.sound] — choice-free. The schema needs no Axiom of Choice.
Number-system instance (ZPO_Ostrowski.lean): [propext, Classical.choice, Quot.sound] — Classical.choice inherited from Mathlib's classical analysis / number theory (Ostrowski).
Core choice-free, realisation choice-carrying — the framework's standing pattern. Zero sorry. Verified: lake build, June 2026.

End of ZP-O v1.0 | The Fixed-Point Fork | fork_collapse_iff (choice-free spine) | Ostrowski number-system instance |
hard fence: cross-instance identity is a type boundary | soft fence: not every fork is μ/ν | All ZPO.lean /
ZPO_Ostrowski.lean theorems verified.